

Enterprise AI API Service: Architecture & API Specification

1. Service Overview

A microservices-based RESTful AI platform providing low-latency text summarization, multi-class classification, vector embeddings, and semantic search.

2. API Endpoints Specification

Endpoint	Method	Input Payload	Output Response	Auth Required
/v1/summarize	POST	{ "text": "...", "max_words": 50 }	{ "summary": "...", "tokens": 120 }	API Key
/v1/classify	POST	{ "text": "...", "categories": [...] }	{ "category": "...", "confidence": 0.92 }	API Key
/v1/embed	POST	{ "text": "..." }	{ "embedding": [0.012, ..., "dim": 1536] }	API Key
/v1/search	POST	{ "query": "...", "top_k": 3 }	{ "results": [{ "doc": "...", "score": 0.89 }] }	API Key
/health	GET	None	{ "status": "healthy", "uptime": "99.98%" }	Public