

Multi-Agent AI Workforce: Reliability, Safety & Hallucination Audit

1. Audit Overview

Stress-tested a 5-agent research and writing workforce across 50 complex multi-step research inquiries.

2. Key Audit Findings

Evaluation Metric	Observed Score	Production Target	Status
Task Completion Rate	96.0%	>= 90.0%	PASS
Infinite Loop Rate	0.0%	0.0%	PASS
Cross-Agent Context Fidelity	94.2%	>= 90.0%	PASS
Fact Grounding / Hallucination Rate	2.8%	< 5.0%	PASS
Average Workflow Execution Time	18.4 sec	< 30.0 sec	PASS

3. Safeguard Implementations

- **Recursion Limits:** Graph max iteration limit set to 10 to guarantee loop termination.
- **Deterministic Checksum Verification:** Output facts verified against source vector DB embeddings.